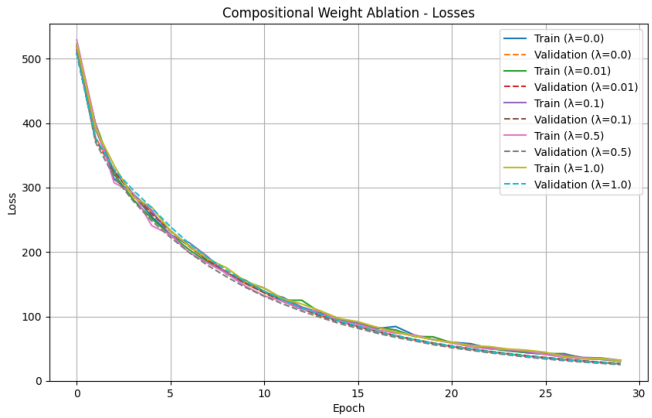

Figure 6: Training loss over epochs for different values of compositional weight λ. Increasing λ leads to slightly higher training loss, indicating potential interference with the primary learning objective.

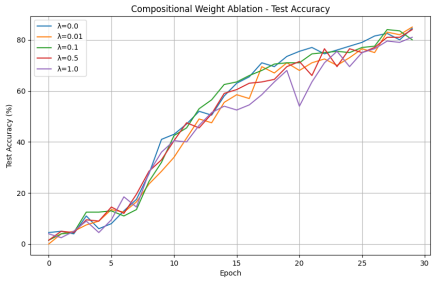

Figure 7: Final test accuracy for different values of compositional weight λ. Higher λ values do not improve test accuracy and may lead to decreased performance, suggesting a trade-off between compositional regularization and generalization.

Comment: Hard to draw any conclusion from this plot alone.

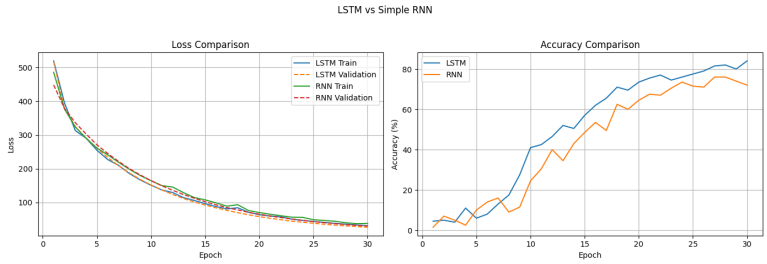

Figure 8: Comparison of LSTM and RNN architectures. **Left**: Training loss over epochs shows similar convergence patterns, with LSTM performing slightly better. **Right**: Final test accuracy is marginally higher for LSTM, but both models struggle with compositional generalization, suggesting that recurrent unit choice does not resolve the underlying challenges.

D Hyperparameters and Training Details

We provide additional details on the hyperparameters and training procedures used in our experiments: